

# Covid 19 analysis

Jie Shen

2023-09-24

## File and Data

This is a R Markdown document for **COVID 19 project for China**. The data used in this project can be found at “[https://github.com/CSSEGISandData/COVID-19/tree/master/csse\\_covid\\_19\\_data/csse\\_covid\\_19\\_time\\_series](https://github.com/CSSEGISandData/COVID-19/tree/master/csse_covid_19_data/csse_covid_19_time_series)”. Please visit the site for detailed data description.

The data I used are global cases and deaths. From the website “<https://github.com/CSSEGISandData/COVID-19>”, Johns Hopkins Corona virus Resource Center ceased its collecting and reporting of global COVID-19 data on March 10, 2023. The global data is from World Health Organization (WHO) “<https://www.who.int/>”.

## Project goal

The project is to discover patterns and trends from COVID data in China. I want to explore things like the COVID cases and deaths trends over the years, and what states are best and worst.

## Packages needed

Be sure the following packages are installed first:

- tidyverse
- ggplot2
- caret
- lubridate

## Load Packages

```
library(tidyverse)
library(ggplot2)
library(forcats)
library(lubridate)
library(dplyr)
library(caret)
```

## Import Data and clean up

```

#Import data from website
url_in<-"https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/csse_covid_19_data/time_series_covid19_confirmed_global.csv"
file_names<-c("time_series_covid19_confirmed_global.csv","time_series_covid19_deaths_global.csv")

urls=str_c(url_in, file_names)
global_cases<-read_csv(urls[1])

```

```

## Rows: 289 Columns: 1147
## -- Column specification -----
## Delimiter: ","
## chr      (2): Province/State, Country/Region
## dbl (1145): Lat, Long, 1/22/20, 1/23/20, 1/24/20, 1/25/20, 1/26/20, 1/27/20,...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

```

```

global_deaths<-read_csv(urls[2])

```

```

## Rows: 289 Columns: 1147
## -- Column specification -----
## Delimiter: ","
## chr      (2): Province/State, Country/Region
## dbl (1145): Lat, Long, 1/22/20, 1/23/20, 1/24/20, 1/25/20, 1/26/20, 1/27/20,...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

```

Now let's take a look and do some clean up

```

# Take a look
head(global_cases)

```

```

## # A tibble: 6 x 1,147
##   'Province/State' 'Country/Region'   Lat   Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>           <dbl> <dbl>   <dbl>   <dbl>   <dbl>
## 1 <NA>            Afghanistan      33.9  67.7         0         0         0
## 2 <NA>            Albania         41.2  20.2         0         0         0
## 3 <NA>            Algeria         28.0   1.66         0         0         0
## 4 <NA>            Andorra         42.5   1.52         0         0         0
## 5 <NA>            Angola          -11.2  17.9         0         0         0
## 6 <NA>            Antarctica      -71.9  23.3         0         0         0
## # i 1,140 more variables: '1/25/20' <dbl>, '1/26/20' <dbl>, '1/27/20' <dbl>,
## #   '1/28/20' <dbl>, '1/29/20' <dbl>, '1/30/20' <dbl>, '1/31/20' <dbl>,
## #   '2/1/20' <dbl>, '2/2/20' <dbl>, '2/3/20' <dbl>, '2/4/20' <dbl>,
## #   '2/5/20' <dbl>, '2/6/20' <dbl>, '2/7/20' <dbl>, '2/8/20' <dbl>,
## #   '2/9/20' <dbl>, '2/10/20' <dbl>, '2/11/20' <dbl>, '2/12/20' <dbl>,
## #   '2/13/20' <dbl>, '2/14/20' <dbl>, '2/15/20' <dbl>, '2/16/20' <dbl>,
## #   '2/17/20' <dbl>, '2/18/20' <dbl>, '2/19/20' <dbl>, '2/20/20' <dbl>, ...

```

```
head(global_deaths)
```

```
## # A tibble: 6 x 1,147
##   'Province/State' 'Country/Region'   Lat   Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>           <dbl> <dbl>    <dbl>    <dbl>    <dbl>
## 1 <NA>            Afghanistan      33.9  67.7      0      0      0
## 2 <NA>            Albania          41.2  20.2      0      0      0
## 3 <NA>            Algeria          28.0   1.66      0      0      0
## 4 <NA>            Andorra          42.5   1.52      0      0      0
## 5 <NA>            Angola          -11.2  17.9      0      0      0
## 6 <NA>            Antarctica      -71.9  23.3      0      0      0
## # i 1,140 more variables: '1/25/20' <dbl>, '1/26/20' <dbl>, '1/27/20' <dbl>,
## #   '1/28/20' <dbl>, '1/29/20' <dbl>, '1/30/20' <dbl>, '1/31/20' <dbl>,
## #   '2/1/20' <dbl>, '2/2/20' <dbl>, '2/3/20' <dbl>, '2/4/20' <dbl>,
## #   '2/5/20' <dbl>, '2/6/20' <dbl>, '2/7/20' <dbl>, '2/8/20' <dbl>,
## #   '2/9/20' <dbl>, '2/10/20' <dbl>, '2/11/20' <dbl>, '2/12/20' <dbl>,
## #   '2/13/20' <dbl>, '2/14/20' <dbl>, '2/15/20' <dbl>, '2/16/20' <dbl>,
## #   '2/17/20' <dbl>, '2/18/20' <dbl>, '2/19/20' <dbl>, '2/20/20' <dbl>, ...
```

```
# Need to pivot dates to rows
global_cases<-global_cases %>%
  pivot_longer(cols= -c("Province/State", "Country/Region", Lat, Long),
               names_to="date",
               values_to="cases")
head(global_cases)
```

```
## # A tibble: 6 x 6
##   'Province/State' 'Country/Region'   Lat   Long date      cases
##   <chr>           <chr>           <dbl> <dbl> <chr>    <dbl>
## 1 <NA>            Afghanistan      33.9  67.7 1/22/20      0
## 2 <NA>            Afghanistan      33.9  67.7 1/23/20      0
## 3 <NA>            Afghanistan      33.9  67.7 1/24/20      0
## 4 <NA>            Afghanistan      33.9  67.7 1/25/20      0
## 5 <NA>            Afghanistan      33.9  67.7 1/26/20      0
## 6 <NA>            Afghanistan      33.9  67.7 1/27/20      0
```

```
# Do similar things to global deaths
global_deaths<-global_deaths %>%
  pivot_longer(cols= -c("Province/State", "Country/Region", Lat, Long),
               names_to="date",
               values_to="deaths")
```

```
# Combine global cases and deaths
global<- global_cases %>%
  full_join(global_deaths) %>%
  mutate(date=mdy(date)) %>%
  rename(Country_Region='Country/Region',
         Province_State ='Province/State')
```

```
## Joining with 'by = join_by('Province/State', 'Country/Region', Lat, Long,
## date)'
```

```
# Summary statistics
summary(global)
```

```
## Province_State      Country_Region      Lat      Long
## Length:330327      Length:330327      Min.   :-71.950      Min.   :-178.12
## Class :character    Class :character    1st Qu.:  3.934      1st Qu.: -42.60
## Mode  :character    Mode  :character    Median : 21.513      Median :  20.94
##                      Mean  : 19.719      Mean  :  22.18
##                      3rd Qu.: 40.464      3rd Qu.:  90.36
##                      Max.   : 71.707      Max.   : 178.06
##                      NA's   :2286        NA's   :2286
##      date      cases      deaths
## Min.   :2020-01-22      Min.   :      0      Min.   :      0
## 1st Qu.:2020-11-02      1st Qu.:     680      1st Qu.:      3
## Median :2021-08-15      Median :   14429      Median :     150
## Mean   :2021-08-15      Mean   :  959384      Mean   :   13380
## 3rd Qu.:2022-05-28      3rd Qu.: 228517      3rd Qu.:   3032
## Max.   :2023-03-09      Max.   :103802702      Max.   :1123836
##
```

We can see the earliest date is 2020-01-22 and the latest is 2023-03-09.

Since it's unfair to compare the numbers from big population state to a small state, I also want to see cases and deaths per populations. I found the population data set on the same github website.

```
# Import population data
uid_lookup_url="https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/UID
uid=read_csv(uid_lookup_url)
```

```
## Rows: 4321 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (7): iso2, iso3, FIPS, Admin2, Province_State, Country_Region, Combined_Key
## dbl (5): UID, code3, Lat, Long_, Population
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
# After looking through the columns, exclude unwanted columns
uid<-uid%>% select(-c(Lat, Long_, Combined_Key, iso2, iso3, code3,Admin2, UID, FIPS) )
```

```
# Add population column to global data
global<-global%>%
  full_join(uid, by=c("Province_State", "Country_Region"))
```

```
# Get China data
CN<-global%>%filter(Country_Region=="China")
head(CN)
```

```
## # A tibble: 6 x 8
## Province_State Country_Region      Lat      Long date      cases deaths Population
##   <chr>          <chr>      <dbl> <dbl> <date>      <dbl> <dbl>      <dbl>
```

```
## 1 Anhui      China      31.8  117. 2020-01-22      1      0 61027171
## 2 Anhui      China      31.8  117. 2020-01-23      9      0 61027171
## 3 Anhui      China      31.8  117. 2020-01-24     15      0 61027171
## 4 Anhui      China      31.8  117. 2020-01-25     39      0 61027171
## 5 Anhui      China      31.8  117. 2020-01-26     60      0 61027171
## 6 Anhui      China      31.8  117. 2020-01-27     70      0 61027171
```

## Missing Values

Check missing value

```
as.data.frame(colSums(CN%>%(is.na)))
```

```
##               colSums(CN %>% (is.na))
## Province_State                1
## Country_Region                0
## Lat                          1144
## Long                         1144
## date                          1
## cases                         1
## deaths                       1
## Population                    1143
```

clean up missing values

```
# I think we don't need Lat and Long columns, exclude them
CN<-CN%>%select(-c(Lat,Long))
```

```
# Since I will aggregate data by dates, I will just exclude rows with missing date. My guess is those a
CN <- CN %>% filter(!is.na(date))
```

```
# Check what Province_State those rows with missing Population comes from
unique(CN %>% filter(is.na(Population)) %>% filter(!is.na(Province_State))) %>% distinct(Province_State)
```

```
## # A tibble: 1 x 1
##   Province_State
##   <chr>
## 1 Unknown
```

```
# It seems those are all "unknow" states. I want to exclude them as well as cases 0 rows
CN <- CN %>% filter(cases>0 & Population >0)
```

```
# Take a look at missing values again
as.data.frame(colSums(CN%>%(is.na)))
```

```
##               colSums(CN %>% (is.na))
## Province_State                0
## Country_Region                0
## date                          0
## cases                         0
## deaths                       0
## Population                    0
```

## Analysis

Get per state and total Country numbers

```
# China by state total cases, deaths, and death per million population
CN_by_state<-CN%>%
  group_by( Country_Region,Province_State, date) %>%
  summarise(cases=sum(cases), deaths=sum(deaths), Population = sum(Population)) %>%
  mutate(death_per_mill = deaths/Population*1000000) %>%
  ungroup()
```

## 'summarise()' has grouped output by 'Country\_Region', 'Province\_State'. You can  
## override using the '.groups' argument.

```
#Take a look
tail(CN_by_state)
```

```
## # A tibble: 6 x 7
##   Country_Region Province_State date      cases deaths Population
##   <chr>          <chr>      <date>    <dbl>  <dbl>      <dbl>
## 1 China          Zhejiang  2023-03-04 11848      1    64567588
## 2 China          Zhejiang  2023-03-05 11848      1    64567588
## 3 China          Zhejiang  2023-03-06 11848      1    64567588
## 4 China          Zhejiang  2023-03-07 11848      1    64567588
## 5 China          Zhejiang  2023-03-08 11848      1    64567588
## 6 China          Zhejiang  2023-03-09 11848      1    64567588
## # i 1 more variable: death_per_mill <dbl>
```

```
# China Totals
CN_totals<- CN%>%
  group_by( Country_Region, date) %>%
  summarise(cases=sum(cases), deaths=sum(deaths), Population = sum(Population)) %>%
  mutate(death_per_mill = deaths/Population*1000000) %>%
  arrange(death_per_mill) %>%
  ungroup()
```

## 'summarise()' has grouped output by 'Country\_Region'. You can override using  
## the '.groups' argument.


```
#Take a look
tail(CN_totals)
```

```
## # A tibble: 6 x 6
##   Country_Region date      cases deaths Population death_per_mill
##   <chr>          <date>    <dbl>  <dbl>      <dbl>      <dbl>
## 1 China          2023-03-04 3381708 18858 1417925054      13.3
## 2 China          2023-03-05 3381708 18859 1417925054      13.3
## 3 China          2023-03-06 3381708 18860 1417925054      13.3
## 4 China          2023-03-07 3381708 18860 1417925054      13.3
## 5 China          2023-03-08 3381708 18860 1417925054      13.3
## 6 China          2023-03-09 3381708 18861 1417925054      13.3
```

## Visualization CN totals

```
# Visualize CN totals
options(repr.plot.width=30, repr.plot.height=10)
CN_totals %>%
  filter(cases>0) %>%
  ggplot(aes(x=date, y=cases)) +
  geom_line(aes(color="cases")) +
  geom_point(aes(color="cases")) +
  geom_line(aes(y=deaths, color="deaths")) +
  geom_point(aes(y=deaths, color="deaths")) +
  scale_y_log10() +
  scale_x_date(date_labels = "%y-%b", date_breaks = "4 month") +
  theme(legend.position='bottom', axis.text=element_text(angle=90, size=10)) +
  labs(title="COVID 19 in China - total cases and deaths", y=NULL)
```

### COVID 19 in China – total cases and deaths



## How about new cases and new deaths?

When looking at trends, it's good to see how many new cases and new deaths. Let's add those columns

```
# Add new cases columns to China data
CN_by_state<- CN_by_state%>% arrange(Country_Region, Province_State, date) %>%
  mutate(new_cases=cases-lag(cases), new_deaths=deaths-lag(deaths))
```

```
CN_totals<- CN_totals%>% arrange(Country_Region, date) %>%
  mutate(new_cases=cases-lag(cases), new_deaths=deaths-lag(deaths))

# Take a look
tail(CN_by_state)
```

```
## # A tibble: 6 x 9
##   Country_Region Province_State date      cases deaths Population
##   <chr>           <chr>         <date>    <dbl>  <dbl>    <dbl>
## 1 China          Zhejiang      2023-03-04 11848    1   64567588
## 2 China          Zhejiang      2023-03-05 11848    1   64567588
## 3 China          Zhejiang      2023-03-06 11848    1   64567588
## 4 China          Zhejiang      2023-03-07 11848    1   64567588
## 5 China          Zhejiang      2023-03-08 11848    1   64567588
## 6 China          Zhejiang      2023-03-09 11848    1   64567588
## # i 3 more variables: death_per_mill <dbl>, new_cases <dbl>, new_deaths <dbl>
```

```
tail(CN_totals)
```

```
## # A tibble: 6 x 8
##   Country_Region date      cases deaths Population death_per_mill new_cases
##   <chr>           <date>    <dbl>  <dbl>    <dbl>        <dbl>    <dbl>
## 1 China          2023-03-04 3381708 18858 1417925054      13.3        0
## 2 China          2023-03-05 3381708 18859 1417925054      13.3        0
## 3 China          2023-03-06 3381708 18860 1417925054      13.3        0
## 4 China          2023-03-07 3381708 18860 1417925054      13.3        0
## 5 China          2023-03-08 3381708 18860 1417925054      13.3        0
## 6 China          2023-03-09 3381708 18861 1417925054      13.3        0
## # i 1 more variable: new_deaths <dbl>
```

## Visualize new cases and deaths in China

```
# Visualize China totals
options(repr.plot.width=30, repr.plot.height=10)
CN_totals %>%
  filter(cases>0) %>%
  ggplot(aes(x=date, y=new_cases)) +
  geom_line(aes(color="new_cases")) +
  geom_point(aes(color="new_cases")) +
  geom_line(aes(y=deaths, color="new_deaths")) +
  geom_point(aes(y=deaths, color="new_deaths")) +
  scale_y_log10() +
  scale_x_date(date_labels = "%y-%b", date_breaks = "4 month") +
  theme(legend.position='bottom', axis.text=element_text(angle=90, size=10)) +
  labs(title="COVID 19 in China - new cases and deaths", y=NULL)
```

```
## Warning in self$trans$transform(x): NaNs produced
```

```
## Warning: Transformation introduced infinite values in continuous y-axis
```



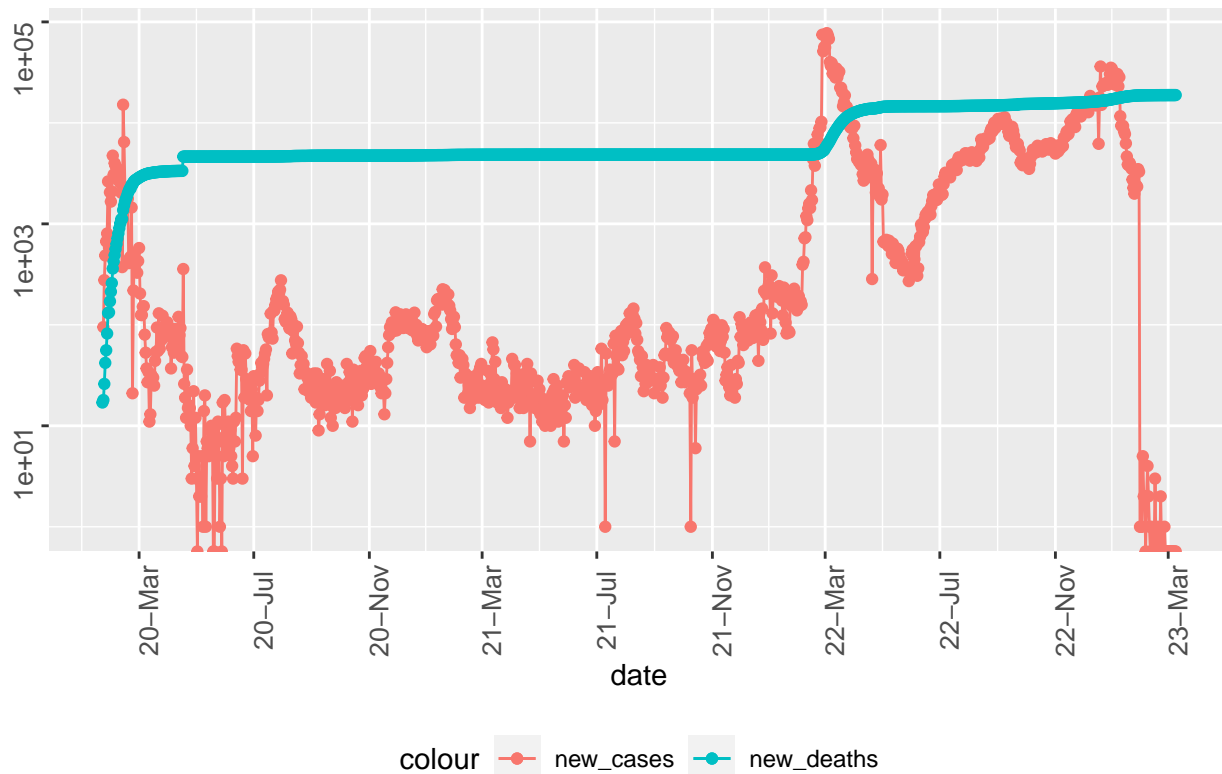
```
## Warning in self$trans$transform(x): NaNs produced

## Warning: Transformation introduced infinite values in continuous y-axis

## Warning: Removed 1 row containing missing values ('geom_line()').

## Warning: Removed 2 rows containing missing values ('geom_point()').
```

## COVID 19 in China – new cases and deaths



## What are the worst and best states in China?

### CN by states

Let's see which states are best/worst (in term of death/population)

```
CN_state_totals <- CN_by_state %>%
  group_by(Province_State) %>%
  summarize(cases=max(cases),
            deaths= max(deaths),
            Population=max(Population),
            cases_per_thou=1000*cases/Population,
            deaths_per_thou=1000*deaths/Population)
CN_state_totals %>% slice_min(deaths_per_thou,n=10)
```

```
## # A tibble: 10 x 6
##   Province_State cases deaths Population cases_per_thou deaths_per_thou
##   <chr>          <dbl> <dbl>    <dbl>          <dbl>          <dbl>
## 1 Jiangsu        5075     0  84748016         0.0599           0
## 2 Ningxia        1276     0   7202654         0.177           0
## 3 Qinghai         782     0   5923957         0.132           0
## 4 Tibet          1647     0   3648100         0.451           0
## 5 Zhejiang      11848     1  64567588         0.183      0.0000155
## 6 Shanxi         7167     1  34915616         0.205      0.0000286
## 7 Guangxi       13371     2  50126804         0.267      0.0000399
## 8 Inner Mongolia  8847     1  24049155         0.368      0.0000416
## 9 Jiangxi        3423     2  45188635         0.0757     0.0000443
## 10 Liaoning      3547     2  42591407         0.0833     0.0000470
```

```
CN_state_totals %>% slice_max(deaths_per_thou,n=10)
```

```
## # A tibble: 10 x 6
##   Province_State cases deaths Population cases_per_thou deaths_per_thou
##   <chr>          <dbl> <dbl>    <dbl>          <dbl>          <dbl>
## 1 Hong Kong    2876106 13467   7496988         384.           1.80
## 2 Macau         3547    121   649342          5.46           0.186
## 3 Hubei        72131   4515  57752557         1.25           0.0782
## 4 Shanghai     67040    595  24870895         2.70           0.0239
## 5 Beijing      40774     20  21893095         1.86           0.000914
## 6 Hainan       10483     6   10081232         1.04           0.000595
## 7 Heilongjiang  6603     18  31850088         0.207           0.000565
## 8 Chongqing    14715     11  32054159         0.459           0.000343
## 9 Henan        9948     23  99365519         0.100           0.000231
## 10 Tianjin     4392     3   13866009         0.317           0.000216
```


## visualize state of interest

I want to visualize the top 3 worst states

```
state<- "Hong Kong"
CN_by_state %>%
  filter(Province_State==state) %>%
  ggplot(aes(x=date, y=cases)) +
  geom_line(aes(color="cases")) +
  geom_point(aes(color="cases")) +
  geom_line(aes(y=deaths, color="deaths")) +
  geom_point(aes(y=deaths, color="deaths")) +
  scale_y_log10() +
  scale_x_date(date_labels = "%y-%b", date_breaks = "4 month") +
  theme(legend.position='bottom', axis.text=element_text(angle=90, size=10)) +
  labs(title=str_c("COVID 19 in ", state, " - total cases and deaths"), y=NULL)
```

```
## Warning: Transformation introduced infinite values in continuous y-axis
## Transformation introduced infinite values in continuous y-axis
```


## COVID 19 in Hong Kong – total cases and deaths



```
state<- "Macau"
CN_by_state %>%
  filter(Province_State==state) %>%
  ggplot(aes(x=date, y=cases)) +
  geom_line(aes(color="cases")) +
  geom_point(aes(color="cases")) +
  geom_line(aes(y=deaths, color="deaths")) +
  geom_point(aes(y=deaths, color="deaths")) +
  scale_y_log10() +
  scale_x_date(date_labels = "%y-%b", date_breaks = "4 month") +
  theme(legend.position='bottom', axis.text=element_text(angle=90, size=10)) +
  labs(title=str_c("COVID 19 in ", state, " - total cases and deaths"), y=NULL)
```


```
## Warning: Transformation introduced infinite values in continuous y-axis
## Transformation introduced infinite values in continuous y-axis
```

## COVID 19 in Macau – total cases and deaths



```
state<- "Hubei"
CN_by_state %>%
  filter(Province_State==state) %>%
  ggplot(aes(x=date, y=cases)) +
  geom_line(aes(color="cases")) +
  geom_point(aes(color="cases")) +
  geom_line(aes(y=deaths, color="deaths")) +
  geom_point(aes(y=deaths, color="deaths")) +
  scale_y_log10() +
  scale_x_date(date_labels = "%y-%b", date_breaks = "4 month") +
  theme(legend.position='bottom', axis.text=element_text(angle=90, size=10)) +
  labs(title=str_c("COVID 19 in ", state, " - total cases and deaths"), y=NULL)
```

## COVID 19 in Hubei – total cases and deaths



## Modeling

Modeling isn't the focus of this project. But I just want to do a couple of quick ones.

### Linear Regression

```
# fit model
mod<-lm(deaths_per_thou ~ cases_per_thou, data = CN_state_totals)
summary(mod)

##
## Call:
## lm(formula = deaths_per_thou ~ cases_per_thou, data = CN_state_totals)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.014039 -0.007515 -0.006868 -0.006497  0.154572
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.252e-03  5.520e-03   1.133   0.266
## cases_per_thou 4.672e-03  8.264e-05  56.531 <2e-16 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.03117 on 31 degrees of freedom
## Multiple R-squared:  0.9904, Adjusted R-squared:  0.9901
## F-statistic: 3196 on 1 and 31 DF,  p-value: < 2.2e-16
```

```
# add predicted results column
```

```
CN_tot_w_pred <- CN_state_totals %>% mutate(pred = predict(mod))
head(CN_tot_w_pred)
```

```
## # A tibble: 6 x 7
```


```
##   Province_State cases deaths Population cases_per_thou deaths_per_thou    pred
##   <chr>          <dbl> <dbl>      <dbl>          <dbl>          <dbl> <dbl>
## 1 Anhui          2275     7    61027171          0.0373          0.000115 0.00643
## 2 Beijing        40774    20    21893095          1.86           0.000914 0.0150
## 3 Chongqing      14715    11    32054159          0.459           0.000343 0.00840
## 4 Fujian         17122     2    41540086          0.412           0.0000481 0.00818
## 5 Gansu           1742     2    25019831          0.0696          0.0000799 0.00658
## 6 Guangdong     103248    10   126012510          0.819           0.0000794 0.0101
```

```
# visualize results
```

```
CN_tot_w_pred %>% ggplot() +
  geom_point(aes(x = cases_per_thou, y = cases_per_thou), color = "blue") +
  geom_point(aes(x = cases_per_thou, y = pred), color = "red") +
  xlim(0,1)+
  ylim(0,0.1) +
  labs(x = "Cases per Thousand", y = "Deaths per Thousand") # Customize axis labels
```

```
## Warning: Removed 25 rows containing missing values ('geom_point()').
```

```
## Warning: Removed 7 rows containing missing values ('geom_point()').
```



## KNN

```
### fit model
mod <- train(deaths_per_thou ~ cases_per_thou, data = CN_state_totals, method = "knn")
print(mod)
```

```
## k-Nearest Neighbors
##
## 33 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 33, 33, 33, 33, 33, 33, ...
## Resampling results across tuning parameters:
##
##  k  RMSE      Rsquared  MAE
##  5  0.2211850  0.5503169  0.06949801
##  7  0.2225704  0.5425577  0.06942550
##  9  0.2234967  0.5325535  0.06939834
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 5.
```

```
# add predicted results column
```

```
CN_tot_w_pred <- CN_state_totals %>% mutate(pred = predict(mod))
head(CN_tot_w_pred)
```

```
## # A tibble: 6 x 7
```


```
## Province_State cases deaths Population cases_per_thou deaths_per_thou pred
## <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Anhui 2275 7 61027171 0.0373 0.000115 7.18e-5
## 2 Beijing 40774 20 21893095 1.86 0.000914 2.08e-2
## 3 Chongqing 14715 11 32054159 0.459 0.000343 1.30e-4
## 4 Fujian 17122 2 41540086 0.412 0.0000481 1.30e-4
## 5 Gansu 1742 2 25019831 0.0696 0.0000799 5.49e-5
## 6 Guangdong 103248 10 126012510 0.819 0.0000794 2.13e-4
```

```
# visuailze results
```

```
CN_tot_w_pred %>% ggplot() +
  geom_point(aes(x = cases_per_thou, y = cases_per_thou), color = "blue") +
  geom_point(aes(x = cases_per_thou, y = pred), color = "red") +
  xlim(0,1)+
  ylim(0,0.1) +
  labs(x = "Cases per Thousand", y = "Deaths per Thousand") # Customize axis labels)
```

```
## Warning: Removed 25 rows containing missing values ('geom_point()').
```

```
## Warning: Removed 7 rows containing missing values ('geom_point()').
```





## Bias

Bias can come from different sources.

First, I want to consider the quality of data. How is data collected? Are all cases and deaths captured? Why are there some unknown states with unknown populations? Since I excluded those rows with missing population and states, the total population of China won't be accurate.

Second, for modeling, it seems my predictions are far away from the actual. It seems the bias are large there. What causes this? I think maybe over the years the deaths per thousand and cases per thousand changes so much. It won't be good to just use deaths per thousand to predict cases per thousand. Maybe my modeling methods are wrong and my data doesn't fit well with the model assumptions.

Third, when I draw the modeling graphs, I eliminated some outliers. This may give the wrong impression of the range of deaths per thousand and cases per thousand.